

# K Means Clustering

k-means clustering is a method of vector quantization, originally from signal processing, that aims to partition  $n$  observations into  $k$  clusters in which each observation belongs to the cluster with the nearest mean (cluster centers or cluster centroid), serving as a prototype of the cluster. This results in a partitioning of the data space into Voronoi cells .(Too many fancy words mumbo jumbo, right? Don't worry, I will explain everything in a simple way)

- LinkedIn
- YouTube
- github
- Gmail
- discord

## Introduction

### What is K-means Clustering?

K-means is an unsupervised learning algorithm used to solve clustering problems on unlabeled data.

The main idea of K-means clustering is very simple: it takes the input data and tries to divide it into multiple groups (called clusters). But how many groups should there be? That's where the  $K$  comes in.

The  $K$  in K-means represents the number of clusters we want to divide our data into. We manually choose this value and give it to the algorithm.

Once the value of  $k$  is provided, the algorithm creates  $k$  different clusters. Each cluster has a centroid, which represents the center of that cluster.

The algorithm then assigns each data point to the closest centroid. Here, closest means the centroid with the minimum distance from that data point (usually measured using Euclidean distance ).

After assigning all data points, the algorithm recalculates the centroid of each cluster by taking the mean of all the data points inside that cluster. This newly calculated mean becomes the updated centroid .

This process of:

- assigning data points to the closest centroid
- recalculating the mean to update centroids

is repeated again and again until the centroids stop changing (or change very little).

Once this happens, the algorithm stops, and we get the final `clusters`.

So in short, K-means is an iterative algorithm that continuously updates centroids and reassigned data points until a stable clustering is achieved.

## How does K-means Clustering work?

The working process of the K-means algorithm can be broken down into the following steps:

- 1. Initialize centroids** First, we choose `k` initial centroids. This can be done in multiple ways:

- randomly selecting `k` data points from the dataset
- randomly generating `k` points in the feature space (In practice, smarter methods like k-means++ are often used.)

- 2. Assign data points to the closest centroid** Each data point is assigned to the centroid that has the least distance from it.

- 3. Recalculate centroids** For each cluster, we calculate the mean of all assigned data points. This mean becomes the new centroid of that cluster.


- 4. Repeat the process** Steps 2 and 3 are repeated:

- reassign data points
- update centroids

- 5. Stop condition** The algorithm stops when:

- the centroids no longer change, or
- the change is very small, or
- a maximum number of iterations is reached

After stopping, the final `clusters` represent the output of the K-means algorithm.



# How K-Means Clustering Works (Step by Step)

## Step 1: Choose the number of clusters ( $k$ )

Before the algorithm starts, you must decide how many clusters you want.

For example:

- $k = 2 \rightarrow$  split data into 2 groups
- $k = 3 \rightarrow$  split data into 3 groups

At this stage,  $k$  is just a guess. Later, we'll see how the `elbow method` helps choose a good value.

## Step 2: Initialize cluster centroids randomly

Once  $k$  is chosen, the algorithm randomly places  $k$  points in the feature space.

These points are called `centroids`.

Important:

- Centroids are **not real data points**
- They are just coordinates representing the center of a cluster

Because this step is random, different runs can produce different results.

## Step 3: Assign each data point to the nearest centroid

For each data point:

1. Compute the distance to each centroid (usually `Euclidean distance`)
2. Assign the point to the cluster whose centroid is closest

After this step:

- Every data point belongs to **exactly one cluster**
- Clusters are formed based on proximity

## Step 4: Update the centroids

Now that clusters are formed:

- Compute the `mean` of all points in each cluster
- Move the centroid to this new mean position

This is why it's called `K-Means`:

- $K$  clusters
- Centroids are the `mean` of cluster points

## Step 5: Repeat assignment and update steps

Steps 3 and 4 are repeated iteratively:

- Assign points to the nearest centroid
- Recalculate centroids based on current assignments

This continues until:

- centroids no longer move significantly
- or cluster assignments stop changing
- or a maximum number of iterations is reached

At this point, the algorithm is said to have converged.

## What does K-Means optimize? (SSE)

K-Means is not just moving points randomly. It is minimizing a specific objective function called SSE.

SSE measures how compact the clusters are.

Formally:

$$SSE = \sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2$$

Where:

- $C_i$  is the i-th cluster
- $\mu_i$  is the centroid of cluster  $C_i$
- $x$  is a data point

the intuition is:

- Measure distance between each point and its centroid
- Square the distance
- Sum it for all points and all clusters

Lower SSE means:

- tighter clusters
- points are closer to their centroids

Each iteration of K-Means **guarantees that SSE decreases or stays the same.**

Now you might think, why does SSE always decrease?


- Assignment step reduces distance to nearest centroid
- Update step moves centroid to minimize squared distance

This is why K-Means converges.

However:

- it may converge to a **local minimum**
- not always the global best clustering

## Choosing the right $k$ : The Elbow Method



Choosing  $k$  arbitrarily is risky.

The **elbow method** helps estimate a good value.

1. Run K-Means for different values of  $k$  (e.g.  $k = 1$  to  $k = 10$ )
2. For each  $k$ , compute the SSE
3. Plot  $k$  vs SSE

You would quickly notice that the SSE decreases as  $k$  increases and at some point, improvement slows down sharply

That point looks like an **elbow** in the graph.

- Small  $k \rightarrow$  clusters are too broad  $\rightarrow$  high SSE
- Large  $k \rightarrow$  clusters become tighter  $\rightarrow$  lower SSE
- After a certain  $k$ , adding more clusters gives diminishing returns

The elbow point balances:

- simplicity
- compactness

But there are some limitations you have to keep in mind, k means algorithm is,

- Sensitive to initial centroid placement
- Requires choosing  $k$  beforehand
- Works best with spherical, equally sized clusters
- Sensitive to feature scaling

K-Means works by repeatedly assigning points to the nearest centroid and updating centroids to minimize the sum of squared distances, gradually forming compact and well-separated clusters.

For better understanding of `k means clustering` you can watch [this video](#).

Now, let's see how we can make clusters using `sklearn` and use `k means algorithm` to actually cluster the data.

## Implementing K means clustering

As I said earlier, `k means algorithm` doesn't require any labels for the data. It only requires the data itself. But for the purposes of this article, we will use labeled data. And in this case I'll show you how you can make `clustered` datapoints using `sklearn`.

It's pretty simple.

We can just use `make_blobs` function from `scikit-learn` library to make blobs of data.

```
In [11]: from sklearn.datasets import make_blobs  
data = make_blobs(n_samples=200, n_features=2, centers=4, cluster_std=2, r
```

Let me explain what I just did in the code above.

In the `make_blob` function, we can specify the `number of samples(n_samples)`, the `number of features(n_features)`, the `number of clusters(centers)`, the `standard deviation of the clusters(cluster_std)`, and the `random seed(random_state)`.

`cluster_std` is the standard deviation of the clusters. The higher the value, the more spread out the clusters are.

`make_blob` will return a `2D array` that represents the `features` of the `data` and another `array` that represents the `cluster` of each `data` point.

So, we can separate the `data` and the `cluster`,

```
In [12]: features = data[0]  
features.shape
```

```
Out[12]: (200, 2)
```

We have a feature set of `200` `data` points with `2` `features`.

Let's get the `cluster` of each `data` point,

```
In [13]: clusters = data[1]  
clusters.shape
```

```
Out[13]: (200, )
```

The `cluster` of each `data` point is a `1D array` of `200` `data` points.

Now, I'll turn this into a `dataframe` to keep things easy.

Let's visualize the `data` and the `cluster`.

```
In [14]: import pandas as pd

df = pd.DataFrame(features, columns=['feature1', 'feature2'])
df['cluster'] = clusters
df.head()
```


```
Out[14]:   feature1  feature2  cluster
0    8.371428   2.920836      1
1   -7.748688  10.658091      0
2   -8.862821   5.529014      3
3  -11.734496   4.508595      3
4   -8.058357  -5.180905      2
```

looks good. Let's plot the `data` and the `cluster`.

```
In [15]: import seaborn as sns

sns.scatterplot(data=df, x='feature1', y='feature2', hue='cluster', palett
```

```
Out[15]: <Axes: xlabel='feature1', ylabel='feature2'>
```



Well, the cluster looks fairly random. Let's try to cluster the data using `K-means` algorithm.

```
In [16]: from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=4)
```

Now, we fit the data to the `K-means` algorithm.

I've set the `n_clusters` to `4` because I know that there are 4 clusters in the data. I'll show you how to find the number of clusters later in this article

In [17]: `kmeans.fit(features)`

Out[17]:



The model is trained and we can now get the `cluster` centers and the `labels` of the `data`.

In [18]: `pred_clusters = kmeans.cluster_centers_`  
`print(pred_clusters)`


```
labels = kmeans.labels_
print(labels)
```

```
[[ 4.52827469  2.31311912]
 [-8.75330277  7.4646063 ]
 [-6.64425912 -7.05642342]
 [-2.86089586  8.66203102]]
[[0 1 1 1 2 2 0 1 0 2 2 0 0 2 2 2 3 1 2 2 2 2 1 3 1 3 3 2 3 0 2 2 1 1 3 0 1
 0 1 3 2 3 2 2 1 0 0 2 0 3 3 3 1 0 3 3 2 2 3 0 1 0 2 1 1 2 0 3 1 3 3 1 3 2
 0 2 0 3 1 3 3 0 2 1 1 1 3 0 1 2 3 0 0 0 1 3 0 2 3 1 1 3 2 3 0 1 2 2 1 0
 2 3 1 3 1 3 1 3 3 2 0 1 1 0 3 0 0 3 2 2 3 1 1 0 2 2 3 2 0 3 1 0 0 3 0 1
 2 2 3 1 0 1 2 1 1 0 0 0 3 0 0 1 3 2 0 0 2 0 1 3 2 2 0 2 0 3 3 2 3 2 1 1 1
 3 0 0 0 3 3 2 1 1 3 1 0 3 2 0]
```

Now, we visualize if we have correctly clustered the data like we did before.

In [19]: `from matplotlib import pyplot as plt`  
`fig, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(10,6))`  
`ax1.set_title('K Means')`  
`sns.scatterplot(data=df, x='feature1', y='feature2', hue='cluster', palette='`  
`ax2.set_title('Original')`  
`sns.scatterplot(data=df, x='feature1', y='feature2', hue=labels, palette='`

Out[19]: <Axes: title={'center': 'Original'}, xlabel='feature1', ylabel='feature 2'>



the colors don't mean anything. They are just for visualization purposes.

SO, as you can see the clusters are almost the same as before. The `K-means` algorithm has done a good job of clustering the `data`.

Adn sometimes this algorithm will fix the clusters in a different order. But the main thing is that the `K-means` algorithm has clustered the `data` into `4` clusters and it has done a good job of clustering the `data`.

Now let's say we don't know the number of clusters in the data. We can use the `Elbow method` to find the number of clusters in the data.

## Elbow method

The `Elbow method` is a method that is used to find the number of clusters in the data. It is a simple method that is used to find the number of clusters in the data.

As I told you earlier we can use `SSE(Sum of Squared Errors)` to find the number of clusters in the data.

Good for use the `kmeans` after training does give the `SSE` value via the `inertia` attribute.

So, what we can do is select a range of `k` values and then calculate the `SSE` for each value and then plot the `SSE` values against the `k` values.

And find the best `k` for the data.

In [20]: `sse = []`

```
for i in range(1, 11):
```


```

kmeans = KMeans(n_clusters=i, random_state=42)
kmeans.fit(features)
sse.append(kmeans.inertia_)

plt.plot(range(1, 11), sse, marker='o', markerfacecolor='red', linestyle='')

plt.xlabel('Number of Clusters')
plt.ylabel('SSE')
plt.show()

```



And we have a graph of `sse vs k` where we can see that the `sse` for `k=4` is the when it starts to sharply slow down and then it decreases by a small margin and we can select `k=4` as the best `k` for the data.

And that's it for `k means clustering`.

## Final Words

Colors and values mean nothing.