

Natural Language Processing (NLP)

NLP is a fascinating field at the intersection of computer science and linguistics, and it's a key component of many of the technologies we use every day, from search engines to virtual assistants.

In this Article, we'll dive into the core concepts of NLP, explore various techniques, and see how we can apply them to real-world problems. Whether you're a seasoned data scientist or just starting out, there's something here for you.

- [LinkedIn](#)
- [YouTube](#)
- [gtihub](#)
- [Gmail](#)
- [discord](#)

Introduction

Natural Language Processing (NLP) is a subfield of artificial intelligence that focuses on the interaction between computers and humans using natural language. It's a fascinating field that has seen rapid growth in recent years, thanks to advances in machine learning and deep learning.

Some Real World Applications of NLP are:

- Chatbots
- CHAT GPT
- Copilot

Even 30% of this article was written by Copilot 😊

What is NLP

Natural Language Processing (NLP) is a very broad field that encompasses a wide range of tasks, from simple text processing to complex language understanding. It is concerned with the interaction between computers and humans using natural language.

Suppose You work for a customer service department and you receive hundreds of emails every day. It would be impossible to read and respond to each one manually.

Or you are a doctor and you have to go through hundreds of medical records to diagnose a patient. It would be very time-consuming and error-prone for you to do this manually.

There are Hundreds of such examples where if it's done manually it would be very time-consuming and error-prone. This is where **NLP** comes in. It can help you **automate** these tasks and make your life easier.

NLP can help you **extract information from text**, **classify** text into different categories, **summarize text**, **translate** text from one language to another, and much more.

So, think of the first scenario where you receive hundreds of emails every day. You can use **NLP** to **automatically read** and **classify** these emails into different categories. This way, you can **prioritize** which emails to respond to first and which ones to respond to later. How this happens:

- **Compile** all the emails into a single document.
- **Featurize** the text data, meaning you would want to convert the text data into a format that can be used by a machine learning model.
- **Compare** the Features of the text data to a set of predefined categories.

These are the basic steps involved in **NLP** but there are many more advanced techniques that can be used to **extract information** from text data.

How does NLP work?

Here's a simple example to illustrate how **NLP** works:

Suppose you have two **documents** :

- Document 1: "Bob Likes Apples"
- Document 2: "Sam Likes Oranges"

You want to **compare** these two documents to see if they are **similar** or **different**. You can:

- **Tokenize** the documents, meaning we would split the documents into individual words. So, the tokenized version of the documents would be:
 - Document 1: ["Bob", "Likes", "Apples"]
 - Document 2: ["Sam", "Likes", "Oranges"]
- **Vectorize** the documents, meaning we would convert the words into numbers. We can use a technique called **Bag of Words** to do this.

Bag of Words is a simple technique that converts text data into a matrix of word counts. Each row in the matrix represents a document, and each column represents a word. The value in each cell represents the count of the word in the document.

So, we compile all the words in the documents into a single list:

```
[ "Bob", "Sam", "Likes", "Apples", "Oranges" ]
```

Now, we can convert the documents into vectors:

- Document 1:

"Bob Likes Apples"

```
-> [{"Bob": 1,  
 "Sam": 0,  
 "Likes": 1,  
 "Apples": 1,  
 "Oranges": 0}]
```

```
-> [1, 0, 1, 1, 0]
```

- Document 2:

"Sam Likes Oranges"


```
-> [{"Bob": 0,  
 "Sam": 1,  
 "Likes": 1,  
 "Apples": 0,  
 "Oranges": 1}]
```

```
-> [0, 1, 1, 0, 1]
```

Now, We have a fully **vectorized** version of each document. We can now **compare** these vectors to see if they are **similar** or **different**. This is very useful for **document classification** because we are treating the documents as **vectors of features**. SO, we can perform **mathematical operations** like **dot products** and **cosine similarity** to compare the documents.

Now, I'm not going to go deep into the **mathematical details** of how these operations work, GO DO YOUR OWN RESEARCH 😊

COSINE SIMILARITY is a the **dot product** of two vectors divided by the **product** of the **magnitude** or **length** of the two vectors from the **origin**.



The **equation** for **cosine similarity** is:

$$\text{Cosine Similarity} = \frac{A \cdot B}{\|A\| \times \|B\|}$$

Where A and B are the two vectors(vectorized documents) and $\|A\|$ and $\|B\|$ are the magnitudes of the two vectors.

I can also re-write the equation as:

$$\text{Cosine Similarity} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n A_i^2} \times \sqrt{\sum_{i=1}^n B_i^2}}$$

We can use cosine similarity to compare the similarity of two documents . If the cosine similarity is close to 1 , then the documents are similar . If the cosine similarity is close to 0 , then the documents are different .

We can also imporve the Bag of Words model by adjusting the word counts based on the frequency of the words in the corpus .(A corpus is a collection of documents)

- TF-IDF (Term Frequency-Inverse Document Frequency) is a technique that does this. It weights the word counts based on the frequency of the words in the corpus .

TF(Term Frequency) is the importance of the term or word in the document . It is calculated as the number of times the term appears in the document . We represent it as:

$$TF(d, t) = \frac{\text{Number of times } t \text{ appears in } d}{\text{Total number of words in } d}$$

t stands for term and d stands for document .

IDF(Inverse Document Frequency) is the importance of the term or word in the corpus meaning all the documents . It also means how rare the term is in the corpus . It is calculated as the logarithm of the total number of documents divided by the number of documents that contain the term . We represent it as:

$$IDF(t) = \log\left(\frac{D}{dt}\right)$$

Where D is the total number of documents and dt is the number of documents that contain the term .

TF-IDF is calculated as the product of TF and IDF . It weights the word counts based on the frequency of the words in the corpus .

$$W(x, y) = TF(x, y) \times IDF(x)$$

TF-IDF is a very powerful technique that can help you extract important information from text data. We do this to get, not just the word counts, but the importance of the words in the document.

This is just a brief overview of how NLP works. There are many more advanced techniques that can be used to extract information from text data.

Natural Language Processing using Python

Now that we have a basic understanding of NLP, let's see how we can use Python to perform NLP tasks. We'll use the Natural Language Toolkit (NLTK) library, which is a popular library for NLP in Python.

We have to install the NLTK library first. We can do this using the following command:

```
conda install nltk # If you are using Anaconda
```

```
pip install nltk # If you are using pip
```

In this Article, I'll show you the workings of NLP using the NLTK library and build a spam filter using. In this process, we'll learn about tokenization, stemming, lemmatization, and TF-IDF.

Let's get started!

NLTK Basics

The Natural Language Toolkit (NLTK) is a popular library for NLP in Python. It provides a wide range of tools and resources for text processing and analysis. I hope you have already installed the NLTK library.

Lets import the NLTK library and download some resources :

```
In [1]: import nltk
```

Before going to the code, let me give you and overview. I'll use `nltk.download_shell()` to show you how to download the resources. You can download the resources and corpora that you need for your NLP tasks.

This method will open a shell where you can download the resources and corpora that you need. You can download the resources by selecting the number of the resource you want to download.

The shell will give you choices like:

- d to download the resource
- q to quit the shell

- `l` to list the resources
- `u` to update the resources

SO, let's download the resources and corpora named stopwords .

In [2]: `# nltk.download_shell()`

Here I have downloaded the stopwords corpora using the `nltk.download_shell()` method. You can download the resources and corpora that you need for your NLP tasks too.

Now, for info I'll use a dataset from UCI Machine Learning Repository named SMS Spam Collection . This dataset contains SMS messages that are labeled as spam or ham (not spam). We'll use this dataset to build a spam filter using NLP .

You can download the dataset from this link (<https://archive.ics.uci.edu/dataset/228/sms+spam+collection>).

And I also have the dataset in my GitHub repository.

Now, let's start making the spam filter using NLP .

What Kind of Data are we dealing with?

So, we have a dataset that contains SMS messages that are labeled as spam or ham (not spam). But we need to explore the data first to see what kind of data we are dealing with and how we can process it.

So, I'll simply read the dataset using `open()` method and print the first few lines just to get a glimpse of the data and I hope we can understand the data better.

In [3]: `# taking a line from the file and adding it to a list`

```
with open('SMSSpamCollection') as f:  
    messages = f.readlines()  
    messages = [line.rstrip() for line in messages]
```

In [4]: `len(messages)`

Out[4]: 5574

So, you can see that the dataset has 5574 SMS messages that are labeled as spam or ham (not spam). Let's see the first few lines of the data to get a glimpse of the data .

In [5]: `messages[0]`

Out[5]: 'ham\tGo until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat...'

Well this does not look good. there is a \t in the data , what does that mean? Let's explore the data further to see what's going on.

I'm dumbing down the code here, so that anyone can understand it.

let's print the first few lines of the data to see what's going on.

```
In [6]: for mess_no, msg in enumerate(messages[:10]):  
    print(mess_no, msg)
```

```
0 ham Go until jurong point, crazy.. Available only in bugis n great worl  
d la e buffet... Cine there got amore wat...  
1 ham Ok lar... Joking wif u oni...  
2 spam Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005.  
Text FA to 87121 to receive entry question(std txt rate)T&C's apply 0845281  
0075over18's  
3 ham U dun say so early hor... U c already then say...  
4 ham Nah I don't think he goes to usf, he lives around here though  
5 spam FreeMsg Hey there darling it's been 3 week's now and no word back!  
I'd like some fun you up for it still? Tb ok! XxX std chgs to send, £1.50 t  
o rcv  
6 ham Even my brother is not like to speak with me. They treat me like ai  
ds patent.  
7 ham As per your request 'Melle Melle (Oru Minnaminunginte Nurungu Vetta  
m)' has been set as your callertune for all Callers. Press *9 to copy your  
friends Callertune  
8 spam WINNER!! As a valued network customer you have been selected to rec  
eivea £900 prize reward! To claim call 09061701461. Claim code KL341. Valid  
12 hours only.  
9 spam Had your mobile 11 months or more? U R entitled to Update to the la  
test colour mobiles with camera for Free! Call The Mobile Update Co FREE on  
08002986030
```

Well that is interesting. The print statement is splitting the data into two columns based on the tab character. SO, now we know what the \t means. It is the tab character that separates the label from the message .

Which means if we make a dataframe using this data , we have to split the data into two columns based on the tab character.

```
In [7]: import numpy as np  
import pandas as pd  
  
data = pd.read_csv('SMSSpamCollection', sep='\t', names=['label', 'message'])  
data.head()
```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

Here we have a `dataframe` that contains two columns named `label` and `message`. The `label` column contains the labels of the SMS messages (spam or ham) and the `message` column contains the text of the SMS messages.

Now, time for further exploration of the data .

```
In [8]: data.describe()
```

```
Out[8]:
```

	label	message
<code>count</code>	5572	5572
<code>unique</code>	2	5169
<code>top</code>	ham	Sorry, I'll call later
<code>freq</code>	4825	30

We can see that the dataset has 5574 SMS messages that are labeled as spam or ham (not spam).

Now, we can see that the `top` message is a ham message `sorry, I'll call later`. This is the most frequent message in the data .

I want to know the top spam message too. It's because I want to see what kind of spam messages are in the data .

```
In [9]: data.groupby('label').describe()
```

```
Out[9]:
```

label	count		unique		message	
	count	unique	top	freq	count	unique
<code>ham</code>	4825	4516	Sorry, I'll call later	30		
<code>spam</code>	747	653	Please call our customer service representativ...	4		

Feature Engineering

A big part of Machine Learning is feature engineering . It's the process of creating new features from the existing features which can help the machine learning model to learn better.

Feature engineering is a very important step in the machine learning pipeline . It can help you improve the performance of your machine learning model and extract important information from the data .

The more domain knowledge you have, the better you can engineer features from the data . You can create new features from the existing features like I can create a new feature from the message column that contains the length of the message .

```
In [10]: data['length'] = data['message'].apply(len)
```

```
In [11]: data.head()
```

```
Out[11]:
```

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

And now we have an extra column named `length` that contains the `length` of the `message`. This is a `new feature` that we have `engineered` from the `existing features`.

Now time to `visualize` the `data` to see if we can `extract` any `important information` from the `data`.


To do that I'll start with importing the famous `plotting library` `matplotlib` and `seaborn` for `data visualization`.

```
In [12]: import matplotlib.pyplot as plt  
import seaborn as sns
```

Now we can do some `data visualization`.

```
In [13]: plt.figure(figsize=(10, 8))  
sns.histplot(data=data, x='length', kde=True)
```

```
Out[13]: <Axes: xlabel='length', ylabel='Count'>
```



I suggest you play with the `bins` and `kde` values to see how the distribution of the `length` of the `messages` changes.

In my `plot` you can see that the distribution of the `length` of the `messages` is right-skewed. This means that most of the `messages` are short and only a few `messages` are long but there is an exception too. The has a bi-modal distribution going on.

Meaning there are text messages that are quite long. Let's see some details about the `long` `messages`.

```
In [14]: data['length'].describe()
```

```
Out[14]: count    5572.000000
mean      80.489950
std       59.942907
min       2.000000
25%      36.000000
50%      62.000000
75%     122.000000
max      910.000000
Name: length, dtype: float64
```

Well, that's interesting. The `longest` `message` in the `data` is `910` characters long, that's a very long `message`, I wonder if it was from a girl, if that was the case, then I feel sorry for the guy 😊

You know what? I really want to see the `longest` `message` in the `data`.

```
In [15]: print(data[data['length'] == 910]['message'].iloc[0])
```

For me the love should start with attraction.i should feel that I need her every time around me.she should be the first thing which comes in my thoughts.I would start the day and end it with her.she should be there every time I dream.love will be then when my every breath has her name.my life should happen around her.my life will be named to her.I would cry for her.will give all my happiness and take all her sorrows.I will be ready to fight with anyone for her.I will be in love when I will be doing the craziest things for her.love will be when I don't have to prove anyone that my girl is the most beautiful lady on the whole planet.I will always be singing praises for her.love will be when I start up making chicken curry and end up making sambar.life will be the most beautiful then.will get every morning and thank god for the day because she is with me.I would like to say a lot..will tell later..

I couldn't have been more wrong. The longest message is a monologue from a random guy about love and how he feels about it. 😊


With that out of the way, Let's see another histogram of the length of the messages based on the label column. Let's see if we can get some insights from the data .

```
In [16]: plt.figure(figsize=(12, 8))

plt.subplot(1, 2, 1)
sns.histplot(data[data['label'] == 'ham'], x='length', kde=True, color='blue')
plt.title('Ham Message Length')

plt.subplot(1, 2, 2)
sns.histplot(data[data['label'] == 'spam'], x='length', kde=True, color='red')
plt.title('Spam Message Length')
```

```
Out[16]: Text(0.5, 1.0, 'Spam Message Length')
```



SO, we can see that the `spam` messages are `longer` than the `ham` messages. `Ham` messages are `short` almost at the `5-100` characters range. But the `spam` messages are long and are in the `120-180` characters range.

Now, time for the next step,

Text Preprocessing

`Text preprocessing` is a very important step in `NLP`. It involves `cleaning` and `transforming` the `text data` into a `format` that can be used by a machine learning model. As we have seen before(i'm referring to my other articles), `machine learning models` cannot work with `text data` directly because they are `mathematical models` that require `numerical input`.

But in our hand we have `string` text data. So, we have to `convert` the `text data` into a `format` that can be used by a machine learning model. This is where `text preprocessing` comes in.

In the theory section, I have talked about the `bag of words` model and the `TF-IDF` technique. These are the `techniques` that we can use to `convert` the `text data` into a `format` that can be used by a `machine learning model`.

Tokenization

As I have said before, `tokenization` is the process of `splitting` the `text data` into `individual words`. SO, let's `split` the `data`.

To do all that we will use the `built-in string methods` in Python. We will also use the `split()` method to split the data into individual words.

Let's me first demonstrate the `string methods` that we are going to use to split the data.

```
In [17]: import string
```

```
In [18]: # a sample message
mess = 'Sample message! Notice: it has punctuation.'
```

I have taken a sample string and I'll remove the punctuation from the `string` by using the `string.punctuation` which will return a `string` containing all the punctuation characters.

```
In [19]: string.punctuation
```

```
Out[19]: '!"#$%&\'()*+, -./:;<=>?@[\\]^_`{|}~'
```

Now We use string comprehension to remove the punctuation from the `string` and then use `join()` method to join the characters into a single `string` again.

```
In [20]: no_punc = [char for char in mess if char not in string.punctuation]
# if the character is not in the string.punctuation then add it to the list

print(f'Original message: {mess}')
print(f'After removing punctuation: {no_punc}')

# joining the list of characters to form a string
no_punc = ''.join(no_punc) # this is the way to join a list of characters

print(f'Complete message: {no_punc}' )
```

```
Original message: Sample message! Notice: it has punctuation.
```

```
After removing punctuation: ['S', 'a', 'm', 'p', 'l', 'e', ' ', 'm', 'e',
's', 's', 'a', 'g', 'e', ' ', 'N', 'o', 't', 'i', 'c', 'e', ' ', 'i', 't',
' ', 'h', 'a', 's', ' ', 'p', 'u', 'n', 'c', 't', 'u', 'a', 't', 'i', 'o',
'n']
```

```
Complete message: Sample message Notice it has punctuation
```

Now we have a string that does not contain any punctuation.

Now, what we can do is split the string into individual words using the `split()` method. This will split the `string` into individual words based on the space` character.

```
In [21]: no_punc.split()
```

```
Out[21]: ['Sample', 'message', 'Notice', 'it', 'has', 'punctuation']
```

Wel, that was easy right??? Wrong, this is just the beginning. We cannot just split the data into individual words and call it a day. We have to clean the data first.

By cleaning the data I mean removing the stopwords, lowercasing the data, and stemming the data.

So, we start by removing the `stopwords` from the `data`.

We downloaded the `stopwords` `corpora` before. Now we can use the `stopwords` to remove the `stopwords` from the `data`.

```
In [22]: from nltk.corpus import stopwords  
stopwords.words('english')[:10]
```

```
Out[22]: ['a', 'about', 'above', 'after', 'again', 'against', 'ain', 'all', 'am', 'an']
```

Here are the first few stopwords in the corpora. `stopwords` are most common words that do not contribute much to the meaning of the text. So, we can remove the `stopwords` from the `data` to clean the `data`.

Now, we make another list that will contain the `words` that are not in the `stopwords` `corpora`.

Note: The words that are in the stopwords corpora are in lowercase, so we have to lowercase every word in the data to compare the words with the stopwords.

```
In [23]: clean_msg = [word for word in no_punc.split() if word.lower() not in stopw  
print(f'Original msg with stopwords: {no_punc.split()}')  
print(f'Clean msg: {clean_msg}')
```

```
Original msg with stopwords: ['Sample', 'message', 'Notice', 'it', 'has',  
'punctuation']  
Clean msg: ['Sample', 'message', 'Notice', 'punctuation']
```

As you can see that "it" and "has" are removed from the `data` because they are `stopwords`.

We won't lowercase the `data` yet. Because we are not vectorizing the `data` yet. SO, we have to lowercase the `data` before vectorizing the `data`.

But now we have to think about how we can apply this cleaning to the whole dataset. We can use the `apply()` method to apply the cleaning to the whole dataset.

```
In [24]: def text_process(msg):  
    """  
    1. Remove punctuation  
    2. Remove stopwords  
    3. Return list of clean text words  
    """  
    no_punc = [char for char in msg if char not in string.punctuation]  
    no_punc = ''.join(no_punc)  
    return [word for word in no_punc.split() if word.lower() not in stopw
```

Now we apply!

```
In [25]: data['message'].head(5).apply(text_process)
```

```
Out[25]: 0    [Go, jurong, point, crazy, Available, bugis, n...
1           [Ok, lar, Joking, wif, u, oni]
2    [Free, entry, 2, wkly, comp, win, FA, Cup, fin...
3           [U, dun, say, early, hor, U, c, already, say]
4    [Nah, dont, think, goes, usf, lives, around, t...
Name: message, dtype: object
```

So, the `apply()` method applies the `cleaning` function to the whole dataset and in the end we have a cleaned dataset.

Now, for the stemming part...

Stemming is the process of reducing the words to their root or base form. For example, the `stem` of the word `running` is `run`. This can help us reduce the dimensionality of the data and improve the performance of the machine learning model.

But for our data set, stemming is not a good idea. Because the stemmed words might not make any sense in the context of the SMS messages and also in the messages we have words like `U`, `r`, `ok`, `lol` etc. which are not real words but short forms of real words. So, stemming will not be a good idea for our data.

We will use the `text_process` function in the `vectorization` step.

Vectorization

Now, we have to vectorize the data using the bag of words model. We will use the `CountVectorizer` class from the `sklearn` library to vectorize the data.

Before diving into the vectorization part, let me give you a brief overview of the `CountVectorizer` class.

The `CountVectorizer` class is a `sklearn` class that converts a collection of text documents into a matrix of token counts. Each row in the matrix represents a Word and each column represents a message. The value in each cell represents the count of the word in the message.

As the model transforms the data into a matrix of word counts and has a column for each word in the data, the dimensionality of the data will be very high. So, it makes a sparse matrix to store the data.

what is a sparse matrix?

A sparse matrix is a matrix where most elements are zero, making it efficient to store and process by only saving the non-zero values, unlike dense matrices where most elements are significant. this technique saves memory and computation time in fields like data science, scientific computing, and machine learning

| you can find a more detailed explanation [here](#)

In [26]: `from sklearn.feature_extraction.text import CountVectorizer`

Now we can fit our messages to the CountVectorizer class and `transform` the messages into a matrix of word counts.

The CountVectorizer class has a lot of parameters that you can tune to improve the performance of the machine learning model but we will use the default parameters for now and pass our cleaner function to the CountVectorizer class as the analyzer parameter. This way the CountVectorizer class will clean the data before vectorizing the data using the cleaner function we built before.

In [27]: `#bag of words transformer
bow_transformer = CountVectorizer(analyzer=text_process).fit(data['message'])
print(f"Total number of vocabularies in the bag of words: {len(bow_transformer.vocabulary_)}")`

Total number of vocabularies in the bag of words: 11425

This might take a while to `transform` the data into a matrix of word counts. So, be patient.

My data is `transformed` into a matrix of word counts and it has 11425 unique words in the bag of words even after cleaning the data.

Let's now explore the `bow_transformer` to see what we have achieved. Let's get the 4th message from the data and transform it into a matrix of word counts.

In [28]: `mess4 = data['message'][3]
print(f'Message 4: {mess4}')`

Message 4: U dun say so early hor... U c already then say...

In [29]: `bow4 = bow_transformer.transform([mess4])
print(bow4)`

<Compressed Sparse Row sparse matrix of dtype 'int64'
with 7 stored elements and shape (1, 11425)>

Coords	Values
(0, 4068)	2
(0, 4629)	1
(0, 5261)	1
(0, 6204)	1
(0, 6222)	1
(0, 7186)	1
(0, 9554)	2

In [30]: `bow4.shape`

Out[30]: (1, 11425)

SO, we have successfully transformed the 4th message into a matrix of word counts. You might get confused by the `numbers` in the `matrix` and the `shape` of the `matrix`. But don't worry, I'll explain it to you.

the `shape` of the matrix is `(1, 11425)`. This means that the `matrix` has `1 row` and `11425 columns`. The row represents the `message` and the `columns` represent the `words` in the `data`.

We see the output in something like this.

```
(0, 4068)    2
(0, 4629)    1
(0, 5261)    1
(0, 6204)    1
(0, 6222)    1
(0, 7186)    1
(0, 9554)    2
```

This is a `sparse matrix` representation of the `message`. The `numbers` in the `matrix` represent the `count` of the `word` in the `message`. For example, the `word` at index `4068` appears `2` times in the `message`.

this index is the index of the word in the transformed data we vectorized before.

So, if you look carefully you can see that there are `7 words` in the `message` and the `word` at `index 4068` appears `2` times in the `message`.

So, let's see if it's true or not.

```
In [31]: bow_transformer.get_feature_names_out()[4068]
```

```
Out[31]: 'U'
```

We can use the `get_feature_names()` method to get the whole list of `words` in the `transformed data`.

There you go..

We got the word at index `4068` is `U` and it appears `2` times in the message.

So, we can be assured that the `CountVectorizer` class has transformed the data into a matrix of word counts successfully.

Well, we have done a lot of work so far. We have `cleaned the data`, `transformed the data` into a `matrix of word counts`, and `vectorized the data` (not complete yet). Now what we can do is `TF-IDF the data`.

TF-IDF

`TF-IDF` (Term Frequency-Inverse Document Frequency) is a `technique` that `weights` the `word counts` based on the `frequency` of the `words` in the

```
corpus . It is a very powerful technique that can help you extract important information from text data .
```

I have already talked about the TF-IDF technique in the theory section. So, I'll not go into the details of the TF-IDF technique. But I'll show you how to TF-IDF the data using the `TfidfTransformer` class from the `sklearn` library.

But before that we need to get the `sparse matrix` of the `word counts` that we got from the `CountVectorizer` class. We can use the `bow_transformer` to `transform` the data into a sparse matrix of word counts.

```
In [32]: msg_bow = bow_transformer.transform(data['message'])
print(f"shape of the matrix: {msg_bow.shape}")
```

```
shape of the matrix: (5572, 11425)
```

And we have the `sparse matrix` of the word counts that we got from the `CountVectorizer` class. The shape of the matrix is `(5574, 11425)`. This means that the `matrix` has `5574 rows` and `11425 columns`.

The `rows` represent the `messages` and the `columns` represent the `words` in the `data`.

As this is a `sparse matrix`, it does not `store` the `zero elements`. So, we have a lot of `non-zero elements` in the `matrix`. Let's see it

```
In [33]: msg_bow.nnz # number of non-zero occurrences
```

```
Out[33]: 50548
```

Well, that's a lot of non-zero elements in the `matrix`, almost `51k` non-zero elements. This is because the messages has a lot of unique words.

We can calculate the `sparsity` of the `matrix` by dividing the `number of non-zero elements` by the `total number of elements` in the `matrix`.

```
In [34]: sparsity = (100.0 * msg_bow.nnz / (msg_bow.shape[0] * msg_bow.shape[1]))
print(f"Sparsity: {sparsity}")
```

```
Sparsity: 0.07940295412668218
```

Sparsity tells us how `sparse` the `matrix` is.

In this case, the `sparsity` of the `matrix` is `0.0794`. This means that the `matrix` has `7.94%` `non-zero elements`. This is very `efficient` in terms of `memory` because it does not store the zero elements.

Now we use the `TfidfTransformer` class to `TF-IDF` the data. We will `fit` the `sparse matrix` of the word counts to the `TfidfTransformer` class and `transform` the `sparse matrix` into a matrix of `TF-IDF` values.

```
In [35]: from sklearn.feature_extraction.text import TfidfTransformer
tfidf = TfidfTransformer().fit(msg_bow)
```

```

# transforming the bag of words
tfidf4 = tfidf.transform(bow4) # passing the bag of words of message 4

print(f"TFIDF of message 4: \n{tfidf4}")

# we can also check the idf of a particular word
print(f"IDF of the word 'university': {tfidf.idf_[bow_transformer.vocabulary['university']]}")


TFIDF of message 4:
<Compressed Sparse Row sparse matrix of dtype 'float64'
 with 7 stored elements and shape (1, 11425)>
    Coords      Values
(0, 4068)    0.4083258993338407
(0, 4629)    0.2661980190608719
(0, 5261)    0.2972995740586873
(0, 6204)    0.2995379972369742
(0, 6222)    0.31872168929491496
(0, 7186)    0.4389365653379858
(0, 9554)    0.5385626262927565

IDF of the word 'university': 8.527076498901426

```

We can see the Tf-Idf values of the 4th message in the data. The Tf-Idf values are normalized and weighted based on the frequency of the words in the corpus.

Let's transform the whole data into a matrix of TF-IDF values.

```
In [36]: msg_tfidf = tfidf.transform(msg_bow)
```

Now, we can use any machine learning model to classify the messages as spam or ham now.

But in this case I'll use Naive Bayes classifier to classify the messages as spam or ham. I'll use the MultinomialNB class from the sklearn library.

```
In [37]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB

X_train, X_test, y_train, y_test = train_test_split(data['message'], data[
```

We've split the data into training and testing sets.

Now we need to clean the data and transform it into a word count matrix using CountVectorizer. Then we'll apply TF-IDF transformation using TfidfTransformer.

Finally, we'll use MultinomialNB to classify the messages as spam or ham. To automate this process, we'll use sklearn's pipeline feature.

The pipeline takes a list of tuples, where each tuple contains the name of the step and the transformer or estimator to be chained in the pipeline. The pipeline will have three steps: **CountVectorizer**, **TfidfTransformer**, and **MultinomialNB**.


```
In [38]: from sklearn.pipeline import Pipeline
pipeline = Pipeline([
    ('bow', CountVectorizer(analyzer=text_process)), # bow is the name of
```

```
('tfidf', TfidfTransformer()), # tfidf is the name of the step and Tfif
('classifier', MultinomialNB()) # classifier is the name of the step a
])
```

Now, we can directly `fit` the `pipeline` to the `data` and `predict` the `output`.

```
In [39]: pipeline.fit(X_train, y_train)
```

Out[39]:



And we have `trained` the `pipeline` on the training data and predicted the `output` on the testing data using the `MultinomialNB` classifier.

Now, we can use the `classification_report` and `confusion_matrix` to evaluate the performance of the `classifier`.

```
In [40]: predictions = pipeline.predict(X_test)
```

```
In [41]: from sklearn.metrics import classification_report, confusion_matrix
```

```
print(classification_report(y_test, predictions))
print('\n')
print(confusion_matrix(y_test, predictions))
```

	precision	recall	f1-score	support
ham	0.96	1.00	0.98	1451
spam	1.00	0.73	0.84	221
accuracy			0.96	1672
macro avg	0.98	0.86	0.91	1672
weighted avg	0.97	0.96	0.96	1672

```
[[1451    0]
 [ 60   161]]
```

Looks pretty good!

We can also use any other machine learning model. Let's say we want to use the `RandomForestClassifier`.

```
In [42]: from sklearn.ensemble import RandomForestClassifier
```

```
rf = Pipeline([
    ('bow', CountVectorizer(analyzer=text_process)),
    ('tfidf', TfidfTransformer()),
    ('classifier', RandomForestClassifier(n_estimators=1000))
```

```
] )  
  
rf.fit(X_train, y_train)  
  
rf_predictions = rf.predict(X_test)
```

```
In [43]: print(classification_report(y_test, rf_predictions, digits=4))  
print('\n')  
print(confusion_matrix(y_test, rf_predictions, ))
```

	precision	recall	f1-score	support
ham	0.9667	1.0000	0.9831	1451
spam	1.0000	0.7738	0.8724	221
accuracy			0.9701	1672
macro avg	0.9833	0.8869	0.9278	1672
weighted avg	0.9711	0.9701	0.9684	1672

```
[[1451    0]  
 [ 50   171]]
```

And the results are in front of you!

Here's a task for you.

Use cross validation and grid search to find the best hyperparameters for the Random Forest classifier of the pipeline.

Final Words

Writing this article took a lot of time and I'm tired af.