

# Introduction to Neural Networks

I thought I won't make a new article for this, Just as a introduction to Neural Networks. Neural networks needs a long mathematical explanation. In this article I'll not go much deep in the mathematics behind this. But just Enough to get you started with Neural Networks and [tensorflow](#).

- [LinkedIn](#)
- [YouTube](#)
- [github](#)
- [Gmail](#)
- [discord](#)

## Introduction

Neural Networks are a set of algorithms , modeled loosely after the human brain , that are designed to recognize patterns . They interpret sensory data through a kind of machine perception , labeling , or clustering of raw input. The patterns they recognize are numerical and contained in vectors , into which all real-world data (images, sound, text, etc.) must be translated .

Without any of the weird words, Neural Networks are a set of algorithms that help us to model complex patterns and prediction problems that are non-linear in nature.

## Perceptron Model

To fully understand neural networks , we need to understand the Perceptron Model .

A perceptron is the most basic building block of an artificial neural network . Invented by Frank Rosenblatt in 1957.

It is a simplified mathematical model of a single biological neuron.

You can think of a Perceptron as a single neuron of a Human Brain . It takes input , processes it, and provides output .

If you think of a Perceptron as a single neuron , this is how it looks like:



It is the simplest form of a neural network .

Amazingly, even back almost 70 years ago, Frank Rosenblatt was able to see the potential of neural networks and stated that...

"The Perceptron may eventually be able to learn, make decisions, and translate languages."

Which is exactly what we are doing today with Deep Learning and Neural Networks and Machine Learning and Artificial Intelligence and in the base of all these, there is the Perceptrons.

And this is why it is important to understand the Perceptron model.

But it was not until the 1980s that Neural Networks started to become popular again. Because in 1969, Marvin Minsky and Seymour Papert published a book called Perceptrons which showed the limitations of Perceptrons. They showed that Perceptrons were limited to linearly separable problems. The book also suggested that Perceptrons were severely limited in what they could learn.

This Statement by Marvin Minsky and Seymour Papert led to the first AI winter where funding and interest in Neural Networks and AI dried up.


SO, back in the 1970s There was not enough computational power to train Neural Networks and Perceptrons were limited to linearly separable problems.

But in the 1980s, Neural Networks started to become popular again with the discovery of the backpropagation algorithm which allowed Neural Networks to learn non-linear functions.

## Perceptrons


Now again think of the screen you are looking at right now. It is made up of pixels. Each pixel has a color and intensity. If you think of a pixel as a feature, then a Perceptron can be used to classify the pixel as black or white. SO, by analyzing the screen we can see, almost always, we are taking multiple inputs and processing them to get output.

So, now suppose we have two inputs ( $X_1, X_2$ ) and we want to classify them as black or white which we will represent as  $y$ . We can represent this as:



Now, the Perceptron will take input from  $X_1$  and  $X_2$ , process it, and provide output as  $y$ .

We can represent the Perceptron as a mathematical model or a function. So, the final perceptron will look like this:




So, now if I say that the  $f(x)$  is a sum then,

$$f(x) = x_1 + x_2$$

Simple as that. This is the most simplistic form of a Perceptron. It takes input, processes it, and provides output.

Now Realistically, we want to be able to adjust some parameters of the Perceptron to make it more flexible and powerful. In the above model there is no flexibility for the Perceptron to learn from the data. So, what we can do is add a weight to the input which will help the Perceptron to learn from the data more efficiently. So, we will add the weights to the input as a multiplication factor. So, the Perceptron will look like this:



So, the Perceptron will take input from  $X_1$  and  $X_2$ , multiply it with weights  $w_1$  and  $w_2$ , process it, and provide output as  $y$ .

$$y = w_1x_1 + w_2x_2$$

So, now we can adjust the weights to make the Perceptron more accurate and select the weights that minimize the error. This will give us more flexibility to the Perceptron to learn from the data.


We can think of the weights as parameters of the Perceptron that we can adjust to make the Perceptron more accurate.

Now here's a brain teaser for you. What if the actual input is  $X_1$  or  $X_2$  is 0? The weights will not have any effect on the output. How do we fix this?

The answer is to add a bias to the Perceptron.

A **bias** is a constant that helps the **Perceptrons** input to have a **non-zero** value.

Bias is a **constant**. That we can add to the input to make sure that the input is never **0**. So, the weights can have an effect on the output. We can represent the **bias** as **b**. And the **Perceptron** will look like this:



So, the **Perceptron** will take **input** from **X1** and **X2**, multiply it with **weights** **w1** and **w2**, add the **bias** **b**, process it, and provide **output** as **y**.


$$y = (w_1x_1 + b_1) + (w_2x_2 + b_2)$$

So, now we can **adjust** the **weights** and **bias** to make the **Perceptrons** more **flexible** even with the **input** as **0**.

We can **generalize** the **Perceptron** as:

$$y = w_1x_1 + w_2x_2 + b$$

Now think of **n** number of **inputs**. We can represent the **Perceptron** as:



Well, here's is the **mathematical model** of the **Perceptron** with **n** number of **inputs** and we can finally generalize the **Perceptron** as:

$$y = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$


$$y = \sum_{i=1}^n w_i x_i + b$$

**b** is the total bias of all the inputs.

With this we can have modeled a biological neuron as a simple perceptron. Later in this article, we will see how we can expand the input be a tensor of information(n-dimensional array) and how we can expand the Perceptron to a Neural Network .

## Artificial Neural Networks

A neural network is a multi-layer perceptron model that is designed to recognize patterns in the data .



A Neural Network is made up of layers of Perceptrons . You can clearly see that there are five layers in the Neural Network above.

- The first layer is called the input layer and it takes the input from the data .
- The middle layers are called the hidden layers and these layers learn to recognize patterns in the data .
- The last layer is called the output layer and it provides the output from the hidden layer .

Every Perceptron in a layer is connected to every Perceptron in the next layer and works as a input for the next layer and so on.

This multi-layer perceptron model is called a Feedforward Neural Network because the input is fed forward through the network and the output is provided at the end .

Every Neural Network consists of three layers :

- Input Layer
- Hidden Layer
- Output Layer

Input Layer: The input layer is the first layer of the Neural Network . It can have n number of neurons/perceptrons where n is the number of features in the data .

**Output Layer:** The output layer is the last layer of the Neural Network. It can have  $m$  number of neurons/perceptrons where  $m$  is the number of classes in the data meaning the number of output we might have. Also if the output is numerical then the output layer is called the regression layer and it can only have one neuron.

**Hidden Layer :** The hidden layer is the middle layer of the Neural Network. A hidden layer can have  $n$  number of layers and it's because of the hidden layer that the Neural Network can learn complex patterns. The hidden layer is the magic of the Neural Network where all the processing is done and the patterns are learned.

Hidden layers are hard to interpret and understand because every Perceptron in the hidden layer takes input from every Perceptron in the previous layer, so the internal workings become more and more mathematically complicated. As the layers get deeper, the internal workings become more and more mathematically complicated.

Hidden layers have high connectivity and are distantly connected to the input and output layers.

A Neural Network becomes a Deep Neural Network when it has more than one hidden layer.

So, the picture above is a Feedforward Neural Network with three layers and three hidden layers. So, we can say that this is a Deep Neural Network.

Previously, we have seen that a Perceptron can be represented as a mathematical model or a function like:

$$y = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

But this is a linear function and if we think of this as a Neural Network, it'll still be a linear function. But a dataset with non linear values cannot actually be represented by a linear function.

So, we need to generalize the Perceptrons to non-linear values.

We can do that with Activation Functions.

## Activation Functions

The  $W$  clearly implies the weight of incoming input.

$b$  is the bias value which is a constant value that helps the Perceptrons input to have a non-zero value. We can also think of the bias as a offset value that makes  $xw$  have to reach a certain threshold before it can have any effect on the output.

For example if  $b = -10$  in  $x \cdot w + b$  then the effects of  $x \cdot w$  won't really start to overcome the bias until  $x \cdot w$  is greater than 10.


After that, the effect will solely be determined by the value of  $w$  or the weight of the input.

That's why it's called a **bias** because it can **bias** the output of the **Perceptron** to be **positive** or **negative** depending on the value of the **bias**.

Now, the **output** of the **Perceptron** is **linear** in nature. It is a **straight line** and it can only **classify** **linearly separable** problems. But in the real world, we have **non-linear** problems that cannot be **classified** by a **linear Perceptron**.

So, we need to **introduce non-linearity** to the **output** of the **Perceptron** to **classify** **non-linear** problems. This is where the **Activation Functions** come in.

To be completely dumb, you can say that **Activation Functions** are functions that are added to the outputs of the **Perceptrons** to transform the output into a **non-linear** form.



There are many **Activation Functions** that are used in **Neural Networks**. Some of the most popular **Activation Functions** are:

- **Sigmoid** (Logistic Function)
- **Tanh** (Hyperbolic Tangent)
- **ReLU** (Rectified Linear Unit)


## Sigmoid Activation Function

The **Sigmoid Activation Function** is a **non-linear** function that is used in **binary classification** problems. It is also called the **Logistic Function**.

The **Sigmoid Activation Function** is represented as:

$$f(x) = \frac{1}{1 + e^{-x}}$$

The Sigmoid Activation Function takes the output of the Perceptron and squashes it between 0 and 1. This is useful in binary classification problems where we need the output to be between 0 and 1. You have done logistic regression before, right? It's the same thing.



## Tanh Activation Function


It is also called the Hyperbolic Tangent.

The Tanh Activation Function is represented as:

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


The Tanh Activation Function takes the output of the Perceptron and squashes it between -1 and 1. This is useful in binary classification problems where we need the output to be between -1 and 1.

## Hyperbolic Tangent (Tanh) Function



## ReLU Activation Function

One the most popular and widely used Activation Function is the ReLU Activation Function . ReLU stands for Rectified Linear Unit .



ReLU is a linear function that returns 0 if the output is less than 0 and returns the output if the output is greater than 0 as shown in the picture above.

The ReLU Activation Function is represented as:

$$f(x) = \max(0, x)$$

This function shows that the `output` will only be `0` if the `output` is less than `0`.

The `ReLU Activation Function` is used in `Deep Neural Networks` because it is `computationally efficient` and it `speeds up the training` of the `Neural Network`. There have been many `improvements` to the `ReLU Activation Function` like the `Leaky ReLU` and the `Parametric ReLU` and there are still research going on to `improve the ReLU Activation Function`.

There are many more `Activation Functions` that are used in `Neural Networks`. But these are the most popular `Activation Functions` that are used in `Deep Neural Networks`.

You can visit the `wiki page of Activation Functions` to learn more about the `Activation Functions` that are used in `Neural Networks`.  
([https://en.wikipedia.org/wiki/Activation\\_function](https://en.wikipedia.org/wiki/Activation_function))

Now the next part is very critical and without any way to visualize it's pretty hard to explain so I'll go thorough them one by one and try my best to explain them also I will attach some links so that you can learn more about them.

## Multi-Class Classification

`Multi-class classification` means the `target` has **more than two classes**.

There are two types:

### Non-Exclusive Classes

A data point can belong to **more than one class**. Example: an image can be both `cat` and `black`.

This is called **multi-label classification**.

### Mutually Exclusive Classes

A data point can belong to **only one class**. Example: an image is either `cat` or `dog`.

This is the standard **multi-class classification** problem.

## Output Layer Design

If the target has `n` classes, the neural network must have `n output neurons`.

The output is a **vector**, where each value represents how strongly the model believes the input belongs to a class.

So instead of:

red, blue, green

we want:

[0.7, 0.2, 0.1]

These are **scores**, which we later convert into **probabilities**.

---

## One-Hot Encoding

Neural networks cannot learn from labels like red, blue, green.

So we convert them into vectors using one-hot encoding.

label	red	blue	green
red	1	0	0
blue	0	1	0
green	0	0	1

Now the model learns **numbers**, not words.

---

## Output Activation

### For Non-Exclusive Classes → Sigmoid

Each output neuron is independent. The model predicts a probability for **each class separately**.

### For Mutually Exclusive Classes → Softmax

Softmax converts raw scores into **normalized probabilities** that:

- are between 0 and 1
- always sum to 1

$$f(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

This forces the network to **choose one class**.

## Cost Function & Learning

The neural network does not know if it is correct. So we use a cost function to measure **how wrong it is**.

It compares:

- $y$  → actual value
- $a$  → predicted value

And returns **one scalar number**: the error.

Training means:

| change the weights to minimize this number.

## Common Cost Functions

### Mean Squared Error (MSE)

$$C = \frac{1}{2n} \sum (y - a)^2$$

Used mostly for regression .

Large mistakes are punished more because of the square.

### Cross-Entropy Loss

For classification:

$$C = - \sum y \log(a)$$

It strongly punishes **confident but wrong predictions**. That is why it works better with Sigmoid and Softmax .

## Gradient Descent

The cost function depends on thousands of weights. So the error surface is not a curve — it is an n-dimensional landscape.

We cannot visualize it. So we walk downhill using derivatives.

That is Gradient Descent .

$$w = w - \alpha \frac{\partial C}{\partial w}$$

- $\alpha$  is the learning rate
- The derivative tells us the direction of steepest increase
- We move in the opposite direction

## Learning Rate

- Too small → slow

- Too large → overshoot and diverge

So the learning rate must be **tuned**.

## Adaptive Optimizers

Instead of using a fixed learning rate, modern optimizers **adapt** it:

- Adagrad
- RMSprop
- Adam

`Adam` combines momentum and adaptive scaling. That's why it is the default in most frameworks.

## Backpropagation (What Actually Happens)

Training has two steps:

### 1. Forward Pass

Input flows through the network. We compute predictions and the `cost`.

### 2. Backward Pass

We compute:

$$\frac{\partial C}{\partial w}$$

for **every weight** using the `chain rule`.

Then we update the weights using `gradient descent`.

This repeats until the cost stops decreasing.

## Final Truth

Neural networks **do not think**.

They:

1. Guess
2. Measure error
3. Adjust weights
4. Repeat

That's it.

Here are some sources for more information:

MATH:

- [Backpropagation Algorithm\(WIKIPEDIA\)](#)

VIDEO:

- [Whats really happening in Backpropagation\(Channel: 3Blue1Brown\)](#)
- [Backpropagation Calculus\(Channel: 3Blue1Brown\)](#)
- [Neural Networks Playlist\(Channel: 3Blue1Brown\)](#)
- [Neural networks from scratch](#)

So, that's it for the `Backpropagation` algorithm. I tries my best to cover all internal concepts of the `Neural Network`. But without proper way to show you guys there's some things that is not possible for me to cover. With that said,

Now, let's move on to the `Neural Network` implementation in `Python` using the `TensorFlow` and `Keras` libraries.

# Neural Network With Python

## Introduction

In this section, we will implement a `Neural Network` in `Python` using the `TensorFlow` and `Keras` libraries. These two are the most popular libraries for `Deep Learning` and `Neural Networks`.

But Why two libraries?

## TensorFlow and Keras

- `TensorFlow` is a `high-level API` for `Deep Learning`.
- `Keras` is a `high-level API` for `Neural Networks`.
- `TensorFlow` is a `low-level API` for `customization` and `control` over the `neural network architecture`.
- `Keras` is a `high-level API` for `quick prototyping` and `experimentation`.

Buuuuuut, these two are now in the same family. Keres was made using `TensorFlow` to implement `Neural Networks` faster.

In 2015 Google released `TensorFlow` and `Keras` in the same year. But these two were separate libraries. Keras used TensorFlow's API to implement `Neural Networks` faster. Then in 2017 google officially adopted `keras` as tensorflow's `high-level` and user friendly `API` for `Neural Networks`.

So, we can do most of the things we need to implement a `Neural Network` in `Python` using `tf.keras`.

## Installation

It's easy to `install` the `TensorFlow` library using the `pip` package manager or `conda` package manager.

You can get more information about the `installation` of `TensorFlow` from the official `TensorFlow` website (<https://www.tensorflow.org/install/pip>)

Here's how you can `install` the `TensorFlow` library using the `pip` package manager:

```
python -m pip install tensorflow
```

In the official website, they said not to use `conda` to install `tensorflow`. So, I will listen to the experts and not install it using `conda`.

For those who has a `nvidia GPU`, you can `install` the library using this command:

```
pip install tensorflow[and-cuda]
```

Remember `tensorflow` is a huuge library... Close to 5GB. So, be sure that u have enough space to install it.

## Introduction to Tensorflow

Think of it as a more powerfull version of `numpy` with `GPU` support. More modular, more functions, more features.

Think of `keras` as `sklearn` but for `Neural Networks`. It has almost everything you need to implement, every type of `Neural Networks` (yes! there are different types).

But for data preprocessing and other things, we can still use `numpy` and `pandas` because `tensorflow` is fully compatible with `numpy` and `pandas`.

Now that we have `installed` the `TensorFlow` library, let's `import` the `TensorFlow` library and test if it is `installed` correctly.

```
In [1]: import tensorflow as tf  
print(tf.__version__)
```

```
2026-02-04 17:24:06.254675: I tensorflow/core/platform/cpu_feature_guard.c:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.  
2.20.0
```

If you installed the `gpu` version run the code below.

```
In [2]: import tensorflow as tf  
  
print(tf.config.list_physical_devices('GPU'))  
  
[PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

For, cpu testing, run the code below.

```
In [3]: tf.reduce_sum(tf.random.normal([1000, 1000]))
```

```
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR  
I0000 00:00:1770204248.207698 501784 gpu_device.cc:2020] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 9653 MB memory: -> device: 0, name: NVIDIA GeForce RTX 3060, pci bus id: 0000:05:00.0, compute capability: 8.6
```

```
Out[3]: <tf.Tensor: shape=(), dtype=float32, numpy=1471.2747802734375>
```

Noice!

I have the `2.20` version and my gpu is `detected`. So, everything is `working` fine.

Now, let's see how good a neural network can be.

## Let's Make A Neural Network

First we need data.

```
In [4]: import numpy as np  
import pandas as pd  
  
df = pd.read_csv("./custom_price_dataset.csv")  
df.head()
```

```
Out[4]:   feature_1  feature_2      price  
0    0.248357    0.699678  55278.192461  
1   -0.059122    0.472327  54599.930210  
2    0.343864    0.049835  54572.806367  
3    0.791545   -0.293438  56222.631937  
4   -0.077037    0.389152  51960.208200
```


This is a `fake_dataset` where you are given prices according to the `features`. SO, you can see that the dataset has three features, `feature1`, `feature2`, and `price`. The `price` is the `target` that we need to `predict` using the `features`.

We can do some visualizations to understand the `data` better.

```
In [5]: import seaborn as sns  
  
sns.set_style('whitegrid')
```

```
sns.set_palette('Dark2')
sns.pairplot(df, height=3, plot_kws={'s': 10})
```

Out[5]: <seaborn.axisgrid.PairGrid at 0x75d26eccb010>



I made custom patterns to make the data more interesting to see if the Neural Network can learn the patterns in the data.

## Training and Testing Data

I suppose that you guys know how to split the data into training and testing data.

If you don't you should go see my other articles and also visit my youtube channel.

In [6]: `from sklearn.model_selection import train_test_split`

In [7]: `X = df.drop(columns='price')`  
`y = df['price']`  
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)`

I'm doing a 80% training and 20% testing split.

```
In [8]: X_train.shape
```

```
Out[8]: (800, 2)
```

```
In [9]: X_test.shape
```

```
Out[9]: (200, 2)
```

If you guys have large data. Then you should scale it before splitting it into training and testing data because the Neural Network is sensitive to the scale of the data .

But for this dataset the features values range is not that large.

```
In [10]: X.describe()
```

```
Out[10]:
```

	feature_1	feature_2
<b>count</b>	1000.000000	1000.000000
<b>mean</b>	5.009666	5.035418
<b>std</b>	2.949330	2.920090
<b>min</b>	-0.826510	-0.813661
<b>25%</b>	2.564901	2.531286
<b>50%</b>	4.932779	5.030279
<b>75%</b>	7.495682	7.538554
<b>max</b>	10.868813	11.148454

## Building the Neural Network

So, let's build the Neural Network .

First, we need to define the layers.

Let's say:

- we want to make a neural network with 3 hidden layers.
  - 1st hidden layer has 4 neurons .
  - 2nd hidden layer has 6 neurons .
  - 3rd hidden layer has 4 neurons .
- Also we need to define the activation function for the hidden layers .  
Let's just say that the activation function for the hidden layers is relu .
- We want to predict the price so the output layer will have 1 neuron and it doesn't need any activation function.

So, how do we do this?

First, we need to import the layers from tensorflow.

As this is a fully connected Neural Network, we need to import the Sequential model from tensorflow.

```
In [11]: from keras.models import Sequential  
# from tensorflow.python.keras.models import Sequential
```

To define the layers we need to import another class from tensorflow, the Dense class.

```
In [12]: from keras.layers import Dense  
# from tensorflow.python.keras.layers import Dense
```

Now we can start making the Neural Network.

First we initialize the model using the Sequential class.

And inside the sequential model we add the layers using the add method.

We can also directly pass a list of layers to the model using the Sequential class.

```
In [13]: model = Sequential()  
model.add(Dense(4, activation='relu')) # hidden layer with 4 perceptrons  
model.add(Dense(6, activation='relu')) # hidden layer with 6 perceptrons  
model.add(Dense(4, activation='relu')) # hidden layer with 4 perceptrons  
model.add(Dense(1)) # output layer
```

Now, we have a fully connected deep Neural Network.

But this can be a little redundant.

So, we can use the second way to make a neural network.

```
In [14]: model : Sequential = Sequential([  
    Dense(4, activation='relu'),  
    Dense(6, activation='relu'),  
    Dense(4, activation='relu'),  
    Dense(1)  
])
```

And we are done... This process looks easier and cleaner.

So, Now for the next step, We have to define how the model will learn.

We have to define the loss function that will calculate the model's performance on the training data.

We need to define the optimizer that will be used to update the weights and biases of the Neural Network.

So, let's define the `loss` function and the `optimizer` now.

We can use a method from the sequential `model` to define the `loss` function and the `optimizer`.

```
In [15]: model.compile(loss='mse', optimizer='adam')
```

Inside the `compile` method, we can pass the `loss` function and the `optimizer` as arguments.

I passed the `loss` function as '`'mse'`' which means `mean_squared_error` and the `optimizer` as '`'adam'`' which means `adam_optimizer`.

You can pass other `loss` functions like '`'mae'`', '`'binary_crossentropy'` and '`'categorical_crossentropy'` and other `optimizers` like '`'sgd'`', '`'rmsprop'` and '`'adam'`'.

And our model is ready to `train`.

Just like other ml algorithms, we can use the `fit` method to `train` the Neural Network.

But The training process is a little different.

We need to pass the `training` data to the `fit` method.

Then we have to `define` the number of `epochs`. What is an `epoch`? It is the number of times the Neural Network will `see` the entire `training` data to `train` the Neural Network and go through the `training` process.

By default the Neural Network will `see` the entire training data `1` time to `train` the Neural Network.

But that is not a good idea. We can pass 100 or 1000 or 10000 `epochs` to the `fit` method to `train` the Neural Network and go through the `training` process.

But in this case 400. And I'll also set the `verbose` parameter to `1` to see the `training` progress of the Neural Network during the `training` process.

```
In [16]: model.fit(x=X_train, y=y_train, epochs=400, verbose=1) # verbose=0, 1, 2,  
Epoch 1/400
```

```
2026-02-04 17:24:09.864579: I external/local_xla/xla/service/service.cc:16  
3] XLA service 0x75d160017400 initialized for platform CUDA (this does not  
guarantee that XLA will be used). Devices:  
2026-02-04 17:24:09.864595: I external/local_xla/xla/service/service.cc:17  
1] StreamExecutor device (0): NVIDIA GeForce RTX 3060, Compute Capability  
8.6  
2026-02-04 17:24:09.881407: I tensorflow/compiler/mlir/tensorflow/utils/dum  
p_mlir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CRAS  
H_REPRODUCER_DIRECTORY` to enable.  
2026-02-04 17:24:09.981320: I external/local_xla/xla/stream_executor/cuda/c  
uda_dnn.cc:473] Loaded cuDNN version 91801  
25/25 ━━━━━━ 1s 1ms/step - loss: 4365114880.0000  
Epoch 2/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4365071360.0000  
Epoch 3/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4365027840.0000  
Epoch 4/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364979200.0000  
Epoch 5/400  
1/25 ━━━━━━ 0s 14ms/step - loss: 4465872384.0000  
I0000 00:00:1770204250.704504 501860 device_compiler.h:196] Compiled clust  
er using XLA! This line is logged at most once for the lifetime of the pro  
cess.  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364938752.0000  
Epoch 6/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364899328.0000  
Epoch 7/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364850176.0000  
Epoch 8/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364763648.0000  
Epoch 9/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 4364537344.0000  
...  
25/25 ━━━━━━ 0s 1ms/step - loss: 5594172.0000  
Epoch 395/400  
25/25 ━━━━━━ 0s 2ms/step - loss: 5592804.0000  
Epoch 396/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 5572678.5000  
Epoch 397/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 5566153.5000  
Epoch 398/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 5561696.5000  
Epoch 399/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 5553607.5000  
Epoch 400/400  
25/25 ━━━━━━ 0s 1ms/step - loss: 5552910.0000
```

Out[16]: <keras.src.callbacks.history.History at 0x75d1c87cf150>

Now, we have a trained neural network. And we can evaluate the model with the testing set. But before that we should verify that the Neural Network is learning from the data in every epoch and minimizing the loss value to make accurate predictions.

How do we see that.

Every neural network has a `history` attribute that stores the `loss` values at each epoch.

So, we can plot the `loss` values at each epoch to visualize the training progress of the Neural Network.

```
In [17]: model.history.history
```

```
Out[17]: {'loss': [4365114880.0,  
 4365071360.0,  
 4365027840.0,  
 4364979200.0,  
 4364938752.0,  
 4364899328.0,  
 4364850176.0,
```

```
...
```

```
 5690763.5,  
 5691020.0,  
 5670999.0,  
 5664703.5,  
 5678179.0,  
 5652080.0,  
 5630784.0,  
 5617936.0,  
 5616102.0,  
 5595054.5,  
 5594172.0,  
 5592804.0,  
 5572678.5,  
 5566153.5,  
 5561696.5,  
 5553607.5,  
 5552910.0]}
```

Let's plot. I'll transform it to a `dataframe` first and use pandas built-in `plot` method.


```
In [18]: loss_df = pd.DataFrame(model.history.history)
```

```
In [19]: loss_df.head()
```

```
Out[19]:      loss  
0  4.365115e+09  
1  4.365071e+09  
2  4.365028e+09  
3  4.364979e+09  
4  4.364939e+09
```

```
In [20]: loss_df.plot()
```

```
Out[20]: <Axes: >
```



This is good. As you can see the loss started to decrease by the 20th epoch and has a steep decline after that and after that it was declining smoothly.

And that raises a question in my mind. Can the `loss` value decrease more. We ran this for 400 epochs. Is it possible?


Thanks to the `tensorflow` we can continue from where we left.

```
In [21]: model.fit(
    X_train,
    y_train,
    epochs=100,
    verbose=1
)
Epoch 1/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5562906.0000
Epoch 2/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5556253.5000
Epoch 3/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5538517.0000
Epoch 4/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5533753.5000
Epoch 5/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5528807.5000
Epoch 6/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5538772.0000
...
Epoch 99/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5327429.0000
Epoch 100/100
25/25 ━━━━━━━━━━ 0s 1ms/step - loss: 5416493.0000
```

```
Out[21]: <keras.src.callbacks.history.History at 0x75d1905ad010>
```

```
In [22]: pd.DataFrame(model.history.history).plot(y='loss', figsize=(12, 6))
```

```
Out[22]: <Axes: >
```



I trained the model 100 more times and the loss value is not decreasing that much and the loss is oscillating.

So, I think 400 epochs is enough.

Sometimes, I just choose a epoch and see the loss graph to find the point where the loss value stops decreasing and retrain it to that epoch for a better model.

Training a model for longer than needed can cause it to overfit and ruin its progress.

Now, that we have a model ready let's evaluate.

## Evaluating the Neural Network

There are a lot of methods to evaluate the performance of the Neural Network . But the most common method is to calculate the loss value of the Neural Network on the testing data .

For loss of the testing data we can just use the evaluate method .

```
In [23]: model.evaluate(X_test, y_test, verbose=0)
```

```
Out[23]: 5999444.0
```

This will return the mean squared error of the Neural Network on the testing data.

Let's see if it's close to the training loss value.

```
In [24]: model.evaluate(X_train, y_train, verbose=0)
```

```
Out[24]: 5336327.0
```

Pretty close. This means the model didn't overfit that much and it's performing well on the testing data.

So, let's see how it's actually performing by predicting the prices of the testing data and finding the mean absolute error and the root mean squared error of the Neural Network on the testing data.

```
In [25]: test_predictions = model.predict(X_test)
```

```
7/7 ━━━━━━━━ 0s 23ms/step
```

```
In [26]: test_predictions[:10]
```

```
Out[26]: array([[67331.5 ],
 [56194.598],
 [70025.57 ],
 [70047.09 ],
 [54981.914],
 [65676.984],
 [73783.88 ],
 [59229.37 ],
 [59358.246],
 [61401.32 ]], dtype=float32)
```

```
In [28]: from sklearn.metrics import mean_absolute_error, mean_squared_error, root_
print(
    f'Mean Absolute Error: {mean_absolute_error(y_test, test_predictions)}'
)

print(
    f'Root Mean Squared Error: {root_mean_squared_error(y_test, test_predictions)}'
)
```

```
Mean Absolute Error: 2068.42349639483
```

```
Root Mean Squared Error: 2449.3759136940803
```

Hey that's not bad.

if we look at the description of the dataframe.

```
In [29]: df.describe()
```

Out[29]:

	feature_1	feature_2	price
<b>count</b>	1000.000000	1000.000000	1000.000000
<b>mean</b>	5.009666	5.035418	65829.736467
<b>std</b>	2.949330	2.920090	6705.261420
<b>min</b>	-0.826510	-0.813661	50000.000000
<b>25%</b>	2.564901	2.531286	60710.004607
<b>50%</b>	4.932779	5.030279	63626.750195
<b>75%</b>	7.495682	7.538554	72866.770999
<b>max</b>	10.868813	11.148454	80000.000000

The minimum and maximum price is 50000 and 80000 respectively.

Our model is predicting prices between 50000 and 80000 with a mean absolute error of 2068 and rmse of 2449 which is a pretty good in the context of the data .

We are working with a very very small neural model.

Before going to the next thing I want to compare it's results with a traditional machine learning model. In this case `Random Forest would be good right?

Let's do that...

```
In [34]: from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor()
rf.fit(X_train, y_train)
```

Out[34]:

▼ RandomForestRegressor ⓘ ⓘ
  
► Parameters

```
In [35]: from sklearn.metrics import mean_absolute_error, mean_squared_error, root_
print(
    f'Mean Absolute Error: {mean_absolute_error(y_test, rf.predict(X_test))}'
)

print(
    f'Root Mean Squared Error: {root_mean_squared_error(y_test, rf.predict(X_test))}'
)
```

Mean Absolute Error: 1311.5986482382473

Root Mean Squared Error: 1665.043243460561

WHAAAAAT!! HOW CAAAAN THIIIIS BEEE??!!! A MEASLY PROBABILISTIC MODEL OUT PERFORMING THE SUPERIOR NEURAL NETWORK. HOOOOWWWWWW!

It is what it is...

Try to find out what could be wrong. Try to tweak some stuff around like the activation functions or the learning rate.

See if you can beat the `random forest` model.

## Saving and Loading Models

Now that we have a trained `Neural Network` we can `save` it using the `save_model` method of the `tensorflow.keras` module and then we can `load` the `model` using the `load_model` function from the `keras.models` module of the `tensorflow` library.

```
In [39]: from keras.saving import save_model  
  
save_model(model, 'first_model.keras')
```

Now you can load the `model` using the `load_model` function from the `keras.models` module.

```
In [43]: from keras.models import load_model
```

```
In [44]: model = load_model('./first_model.keras', compile=False)
```

```
In [47]: model.predict(X_test)
```

7/7 ━━━━━━ 0s 23ms/step

```
Out[47]: array([[67331.5 ,  
                 [56194.598],  
                 [70025.57 ],  
                 [70047.09 ],  
                 [54981.914],  
                 [65676.984],  
                 [73783.88 ],  
                 ...  
                 [59229.37 ],  
                 [59358.246],  
                 [72693.89 ],  
                 [59538.805],  
                 [67625.53 ],  
                 [64433.31 ],  
                 [76328.21 ]], dtype=float32)
```

AAAAANNNND, Voila. We have successfully built and `trained` Our fist `Neural Network` in Python using the `TensorFlow` and `Keras` libraries.

## Final Words

Good Luck! That's all I can say.

